

---

# SPARC: A Security and Privacy Aware Virtual Machine Checkpointing Mechanism

Mikhail I. Gofman, Ruiqi Luo, Ping Yang, Kartik Gopalan

Research funded in part by the National Science Foundation



*State University of New York*

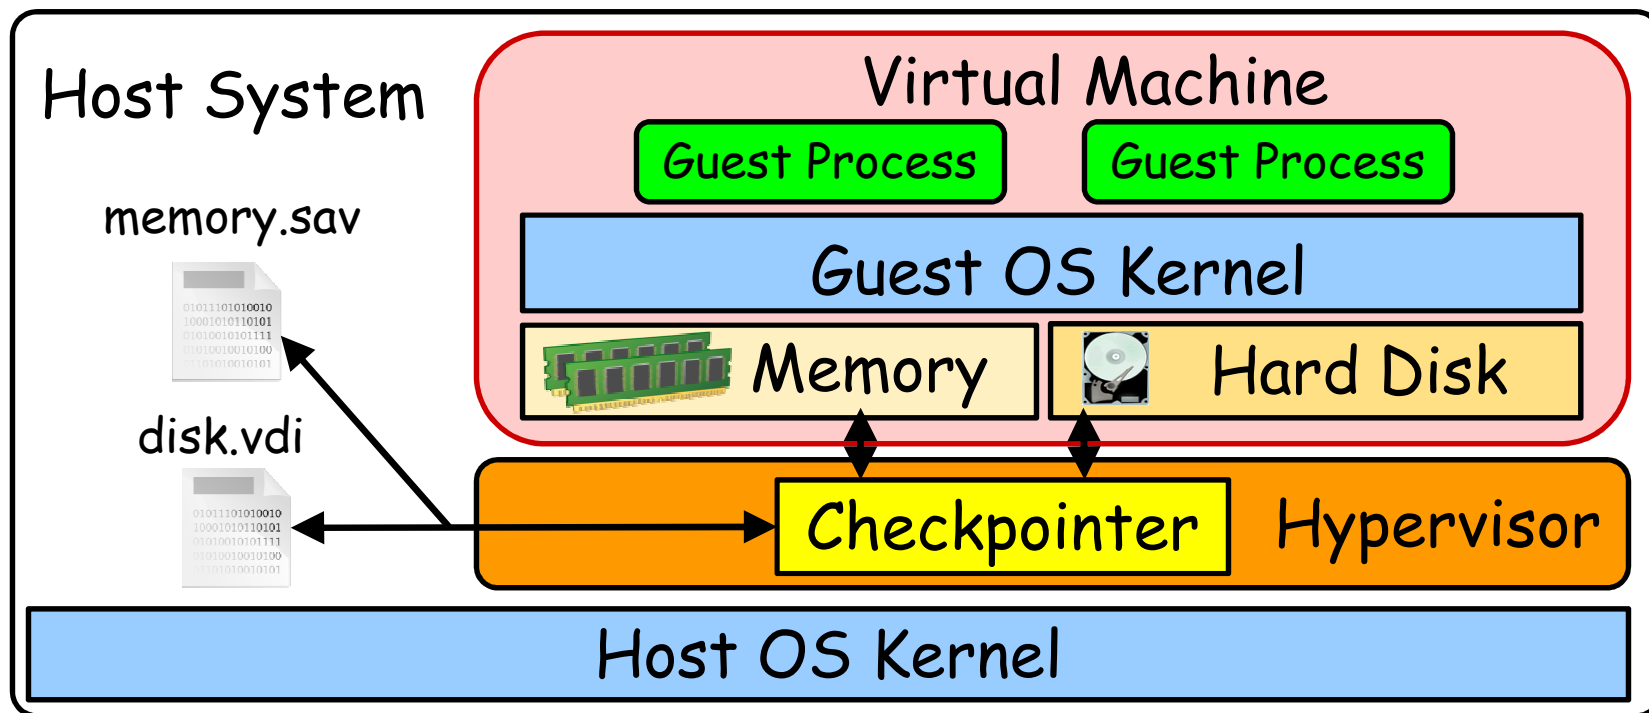
---

# Outline

- Background: Virtual Machine (VM) Checkpointing
- SPARC: a Security and Privacy Aware Checkpointing Mechanism
- Experiments and Performance Results

# Virtual Machine (VM) Checkpointing

- Virtual machine checkpointing saves a snapshot of the physical memory and disk state of a VM in execution:
  - Example: Checkpointing in VirtualBox:
    - Checkpointing:** saves the VM physical memory to a .sav file and disk state to a .vdi file.
    - Restoration:** rolls back the disk state and loads the .sav file into VM memory.



# Virtual Machine (VM) Checkpointing: Benefits and Risks

## ● Benefits:

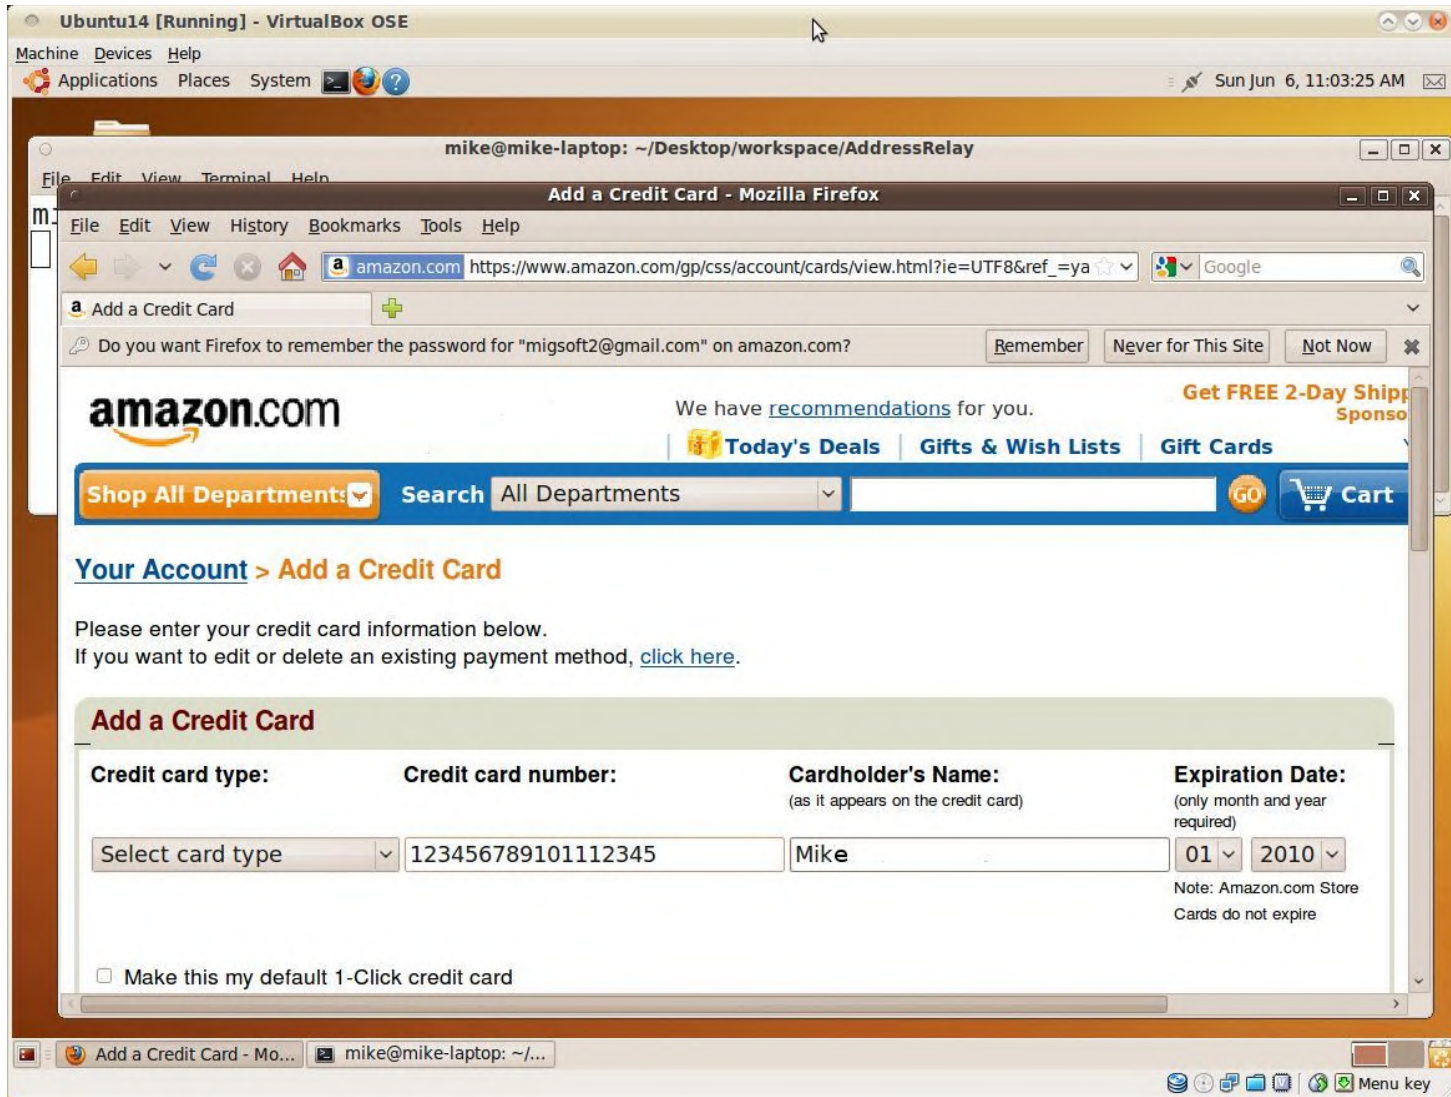
- ◆ Easy recovery of a long running process after it crashes.
- ◆ Easily undo damages caused by malware, patches, etc.

## ● Security Risks:

- ◆ VM memory may contain passwords, credit card #'s, and other sensitive information that should be quickly discarded after usage.
- ◆ Checkpointing drastically prolongs the lifetime of such data by saving it to persistent storage.

# Security Issue: Virtual Machine Checkpointing

- VM checkpoint created/restored using VirtualBox's default checkpointing mechanism:



## Problem Statement

- **Problem:** How can we prevent sensitive data from being leaked via VM checkpoints?
- **Solution:** Prevent sensitive data from being stored in the checkpoint file.

## Existing Approaches

- Clearing deallocated memory:
  - ◆ Does not prevent memory pages from being checkpointed before they are deallocated.
- Protect the checkpointed information by encrypting the checkpoint files:
  - ◆ Restoration decrypts the checkpoint file and loads it into the memory of the VM, thus making the sensitive information vulnerable again.
  - ◆ Attacker can compromise the user's account and gain access to sensitive data by restoring the checkpoints.

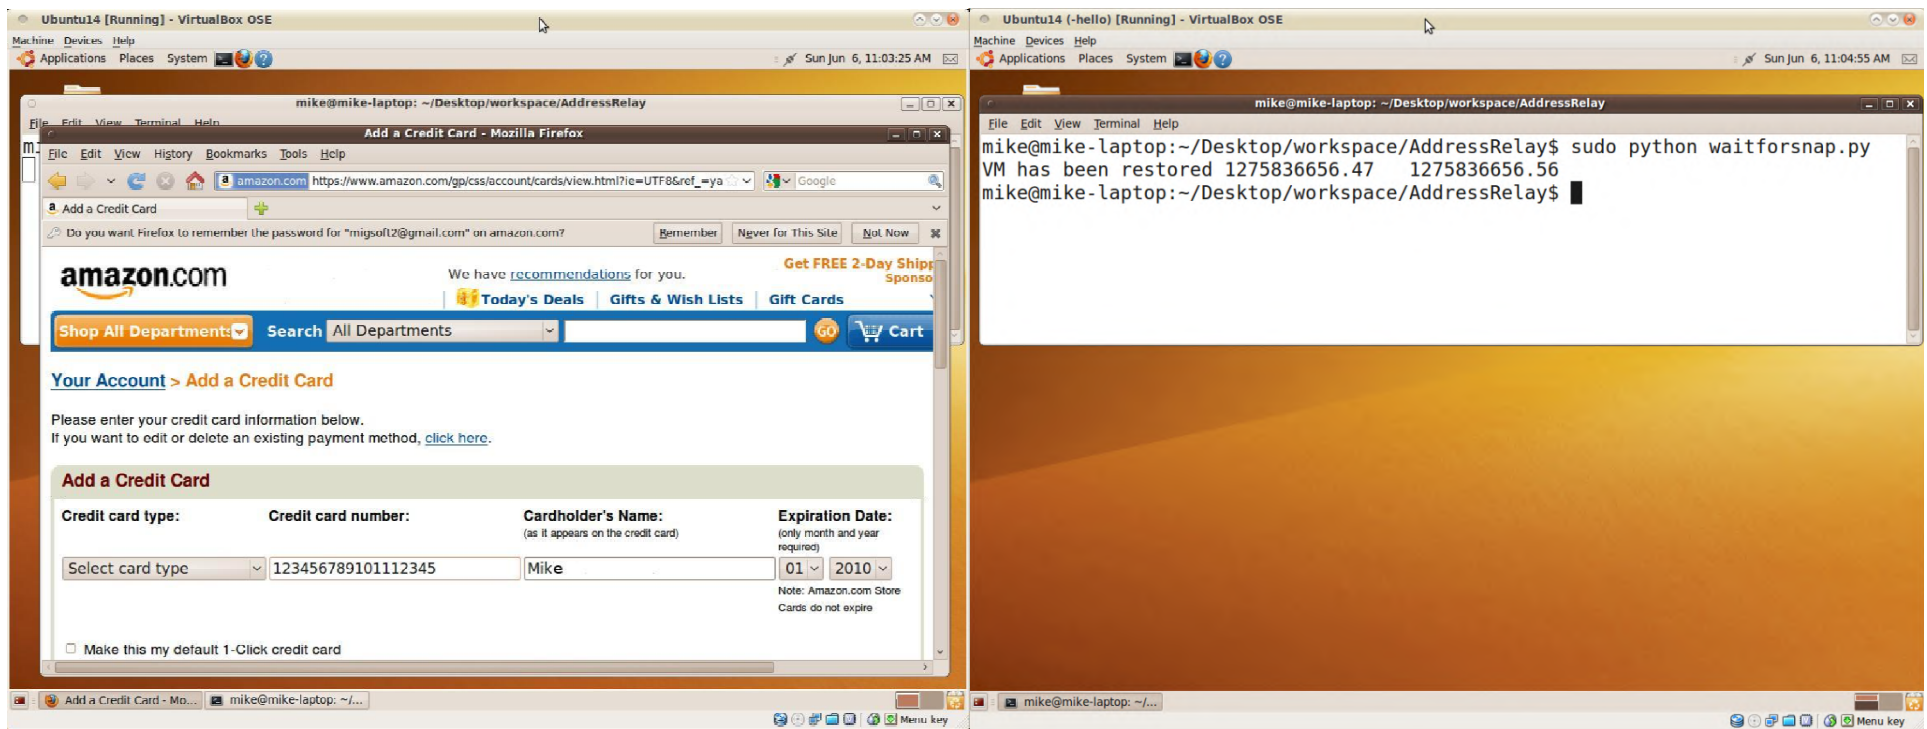
## Contribution

- We developed **SPARC**: a Security and Privacy Aware Checkpointing mechanism:
  - ◆ Enables the users to exclude applications containing user's sensitive information from being checkpointed.



# SPARC in Action

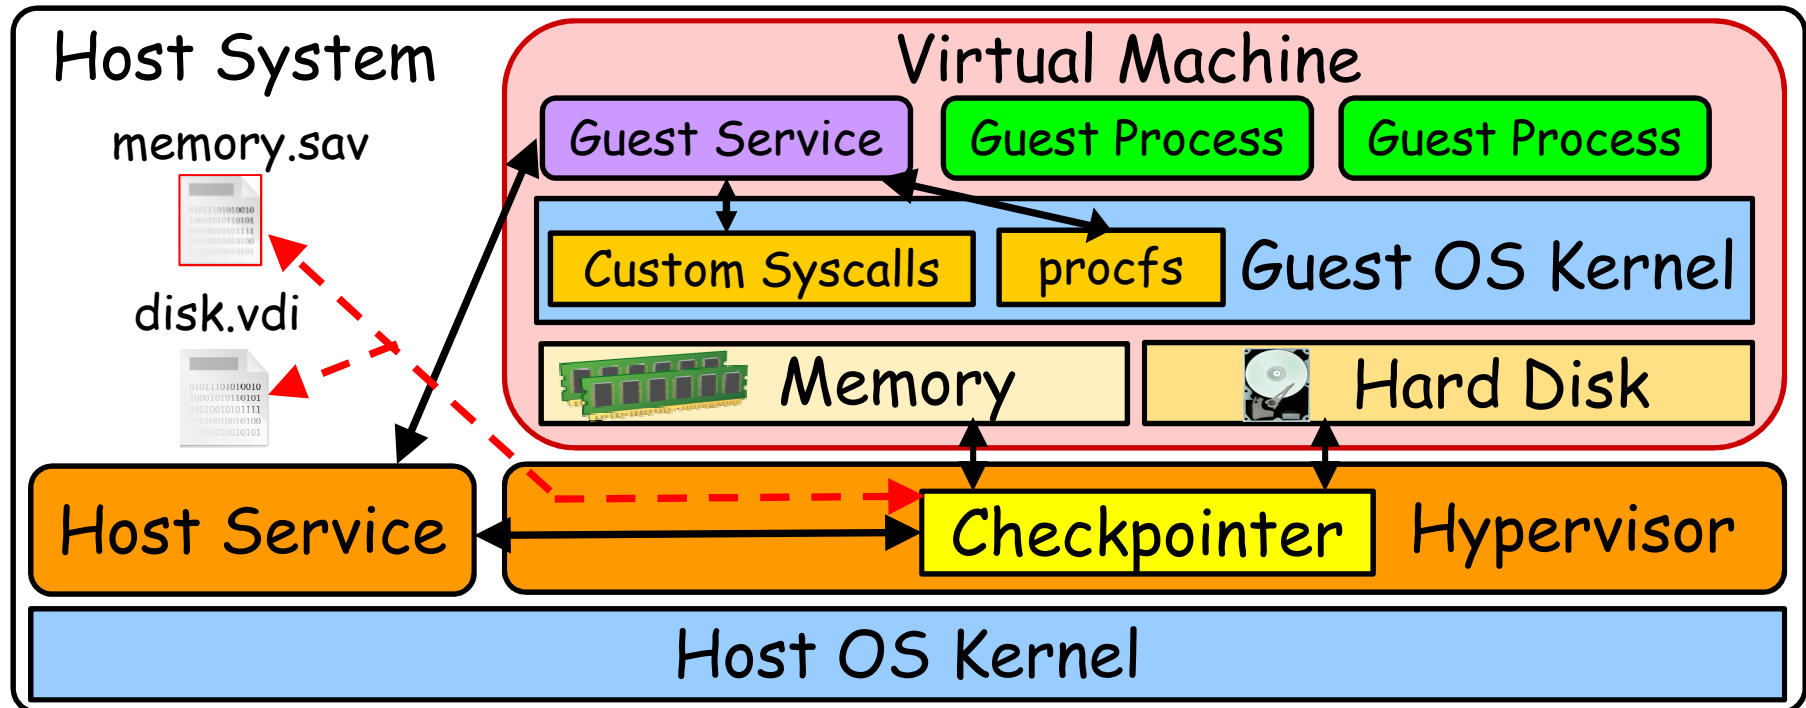
- **Left:** VM checkpoint created/restored using VirtualBox default mechanism.
- **Right:** VM checkpoint created/restored using SPARC with Firefox excluded from the checkpoint.



## SPARC: Key Idea

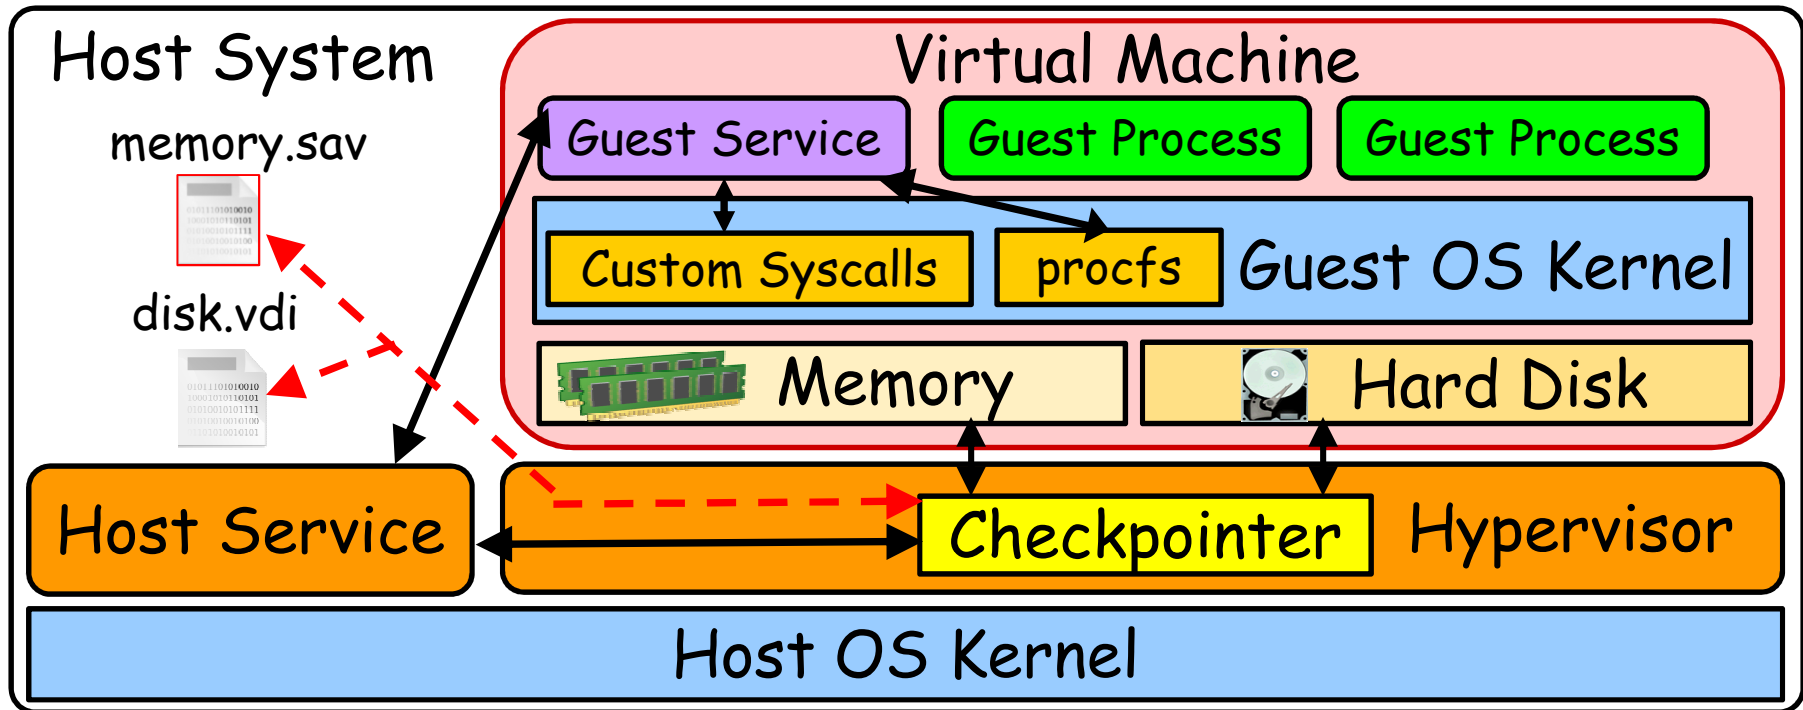
- Track all memory pages containing information related to all privacy sensitive applications and exclude such memory from the checkpoint.

# SPARC Architecture



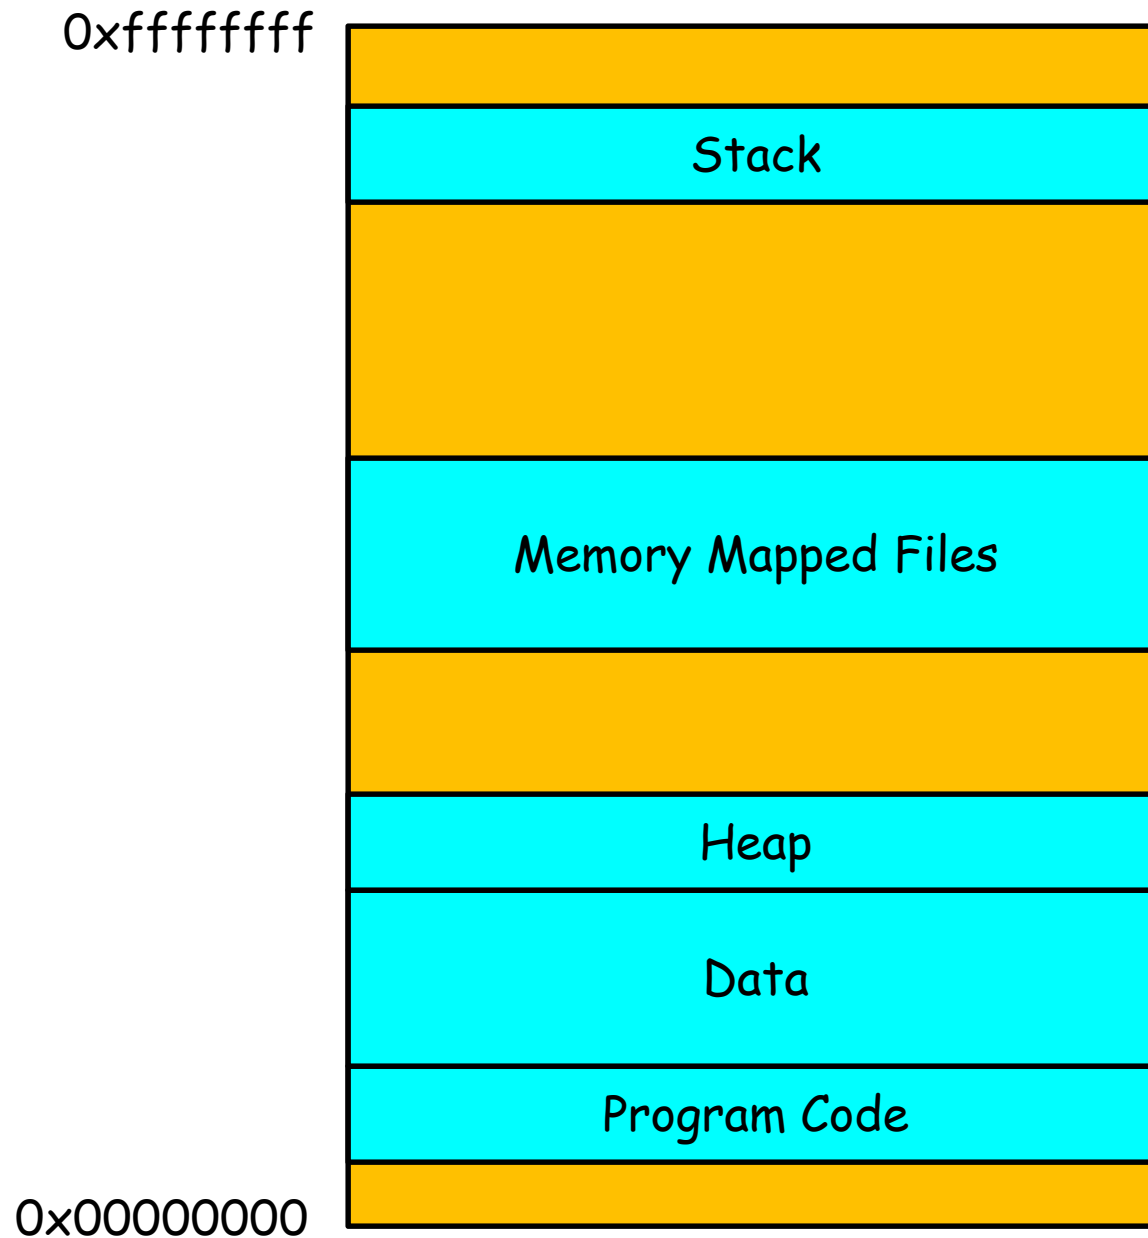
- **Guest Service:** executes inside the VM, detects physical memory addresses of a process to be excluded, and sends them to the host service.
- ◆ **procfs (process file system):** used to identify physical pages belonging to the process.
- ◆ **Custom System Calls:** used to identify kernel physical memory of a process.

## SPARC Architecture (Contd).



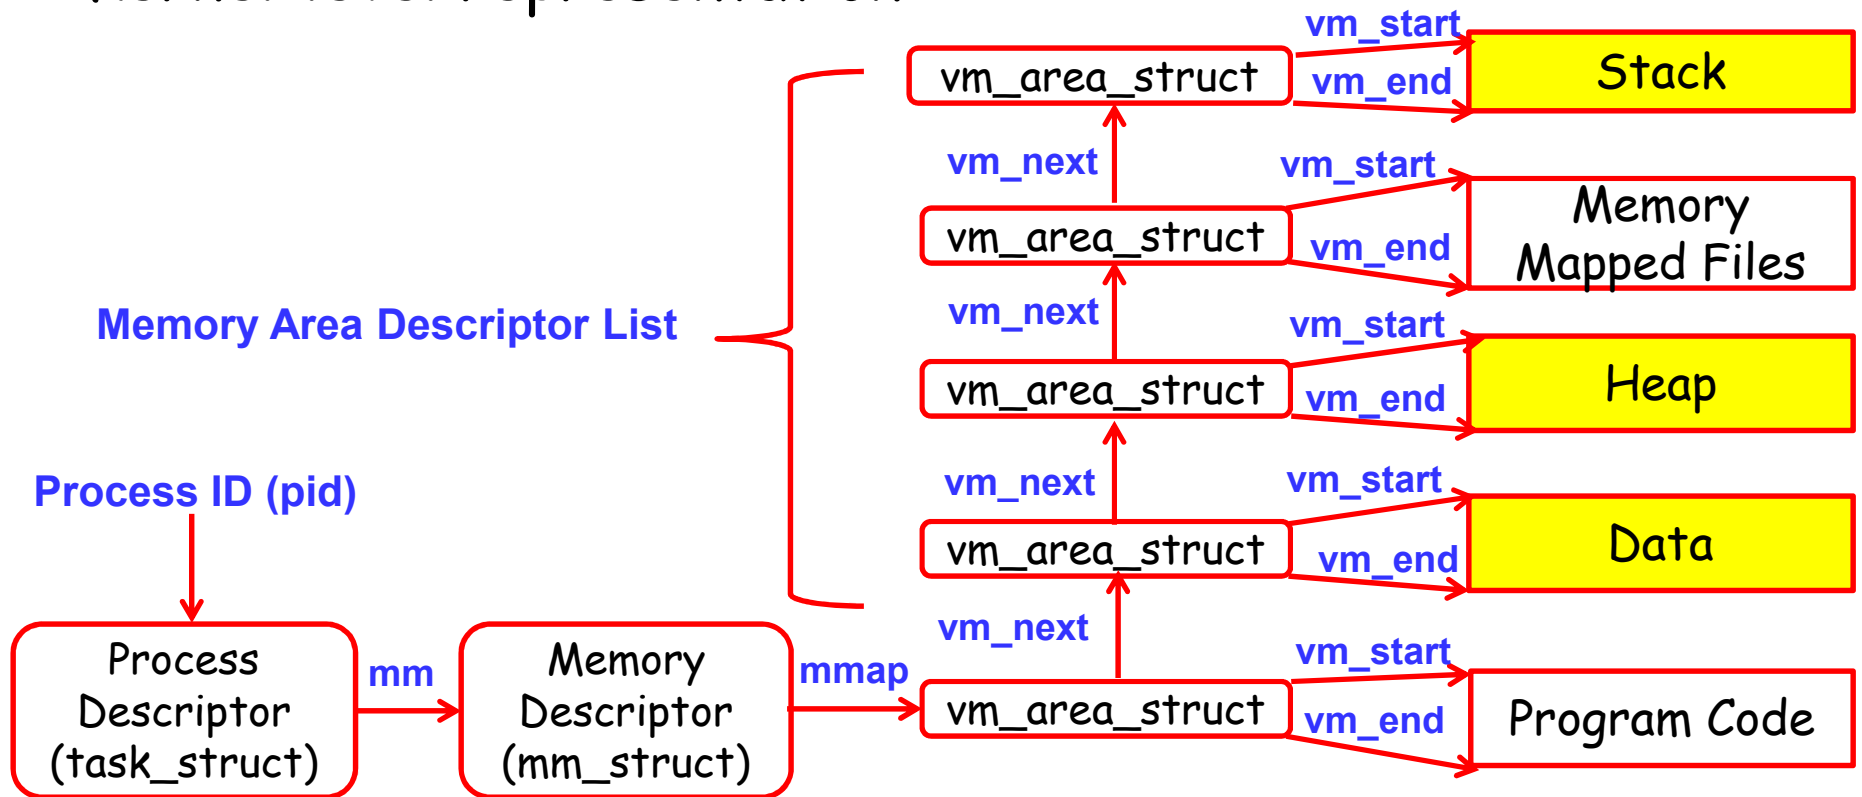
- **Host Service:** receives the physical addresses from the guest service and tells the VirtualBox checkpointer to avoid checkpointing these addresses.

# Process Virtual Address Space



# Identifying Virtual Memory of the Process

- Kernel-level representation:



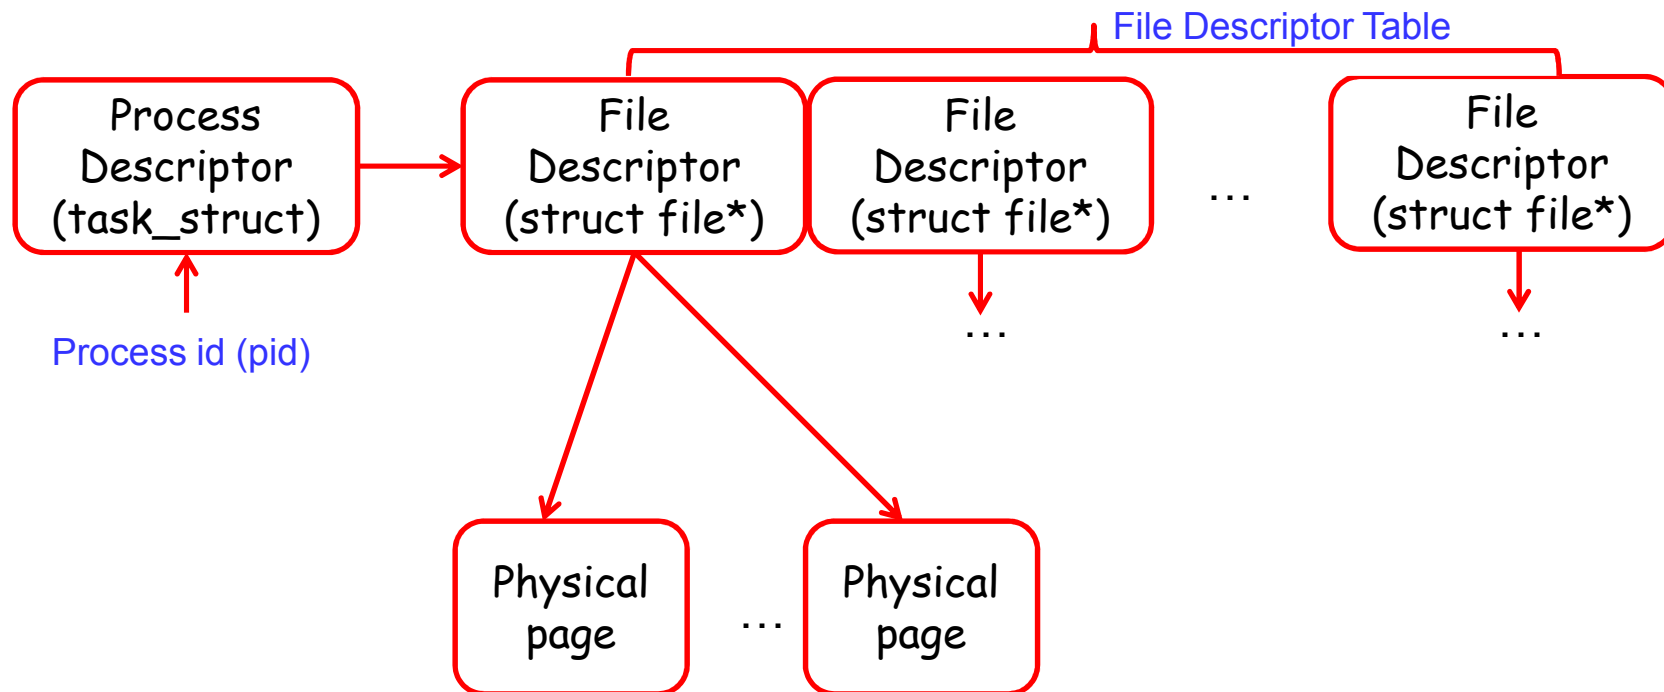
- We traverse the list of memory area descriptors (vm\_area\_structs) and collect starting/ending virtual addresses of **only stack, heap, and data** regions.

# Converting Virtual Addresses to Physical

- Break up virtual memory sections into physical pages and look up corresponding physical addresses in `/proc/pid/pagemap` file.
- **Problem:** virtual-to-physical mappings may change after the physical addresses are collected.
- **Solution:**
  - ◆ Freeze all processes in the VM except the guest service prior to address collection.
  - ◆ Thaw after checkpointing completes and after restoration.
  - ◆ Modify kernel to scrub deallocated memory of the process.

## Excluding pages in the page cache

- May contain **sensitive information** that the process has read/written to/from the file.
  - ◆ Such pages must be excluded from the checkpoint file.





## Excluding pages in the page cache (Contd.)

● **Problem:** After restoration, excluded pages in page cache may affect other processes sharing the same pages.

● **Solution:** Evict all excluded pages after the VM is restored.

● **Problem:** Page cache may retain contents even after a file is closed by a process.

● **Solution:**


- ◆ Evict all pages belonging to such files when the process closes the file.
- ◆ Always clear evicted pages belonging to such files.

## Pipe, FIFO, and Socket Buffers

- Process to be excluded may communicate with other processes through pipes, FIFOs, and sockets.
- Excluding pipe and socket buffers:
  - ◆ Detect file descriptors representing pipes/FIFOs/sockets (similar to detecting descriptors for disk files).
  - ◆ Get the physical addresses from the pipe buffer (`pipe_buffer`) and socket buffer descriptors (`sk_buff`) associated with the file descriptor.

# Excluding Terminal Applications

- Why exclude terminal applications?
  - ◆ Processes may input/output sensitive information to/from the terminals they run on.
- Two types of terminals: Virtual Consoles and Pseudo Terminals
- Terminal applications rely on the Teletype subsystem (TTY) in the kernel.
- Each level of TTY contains buffers which may store sensitive data from the process.



# Excluding Terminal Applications

- Excluding Virtual Consoles:

- ◆ Must detect and exclude all processes running on the console to be excluded:

- Each process descriptor (`task_struct`) contains a pointer to the `tty_struct` representing the associated console.
- Exclude the process and all its descendants.

- Excluding pseudo terminals: similar to excluding Virtual Consoles:

- ◆ **Difference:** must also exclude the associated pseudo terminal driver (pty).

## Handling Restoration

- After restoration we kill the excluded process to allow the OS to clean up any residual state.
- SPARC does not affect the present execution of the VM since the memory is only cleared from the checkpoint file (i.e. not the VM RAM).

# Outline

- Background: Virtual Machine (VM) Checkpointing
- SPARC: a Security and Privacy Aware Checkpointing Mechanism
- Experiments and Performance Results

## Evaluating SPARC's Effectiveness


- Enter a string in xterm terminal, checkpoint the VM, and study the .sav file with a hex editor:
- Without SPARC: string appears 6 times.
- With SPARC:
  - ◆ Excluding the memory of xterm and bash: string appears 3 times.
  - ◆ Excluding the memory of xterm, bash, and the associated TTY subsystem: string disappears.

## Performance Results: Experimental Setup

- We compared the execution times for creating/restoring checkpoints using VirtualBox default checkpointing mechanism and SPARC.
- **Experimental Setup:**
  - ◆ Freshly booted VM running system services, guest service, and a process to be excluded.
  - ◆ Run a program to dirty most of VM's physical pages:
    - VirtualBox reduces checkpoint file size by checkpointing only dirty pages.
  - ◆ Create and restore checkpoint using default mechanism or SPARC and time the operations.
  - ◆ Delete the checkpoint and reboot the VM between subsequent runs.
- Each data point is an average of **five** runs.




# Performance Results: Checkpointing Overhead of SPARC



● SPARC imposes 0.5%–7.1% overhead on checkpointing.

# Performance Results: Restoration Overhead of SPARC



● SPARC imposes 1% - 5.3% overhead on restoration.

# Conclusions and Future Work

## ● Conclusions:

- ◆ VM checkpointing can drastically prolong the lifetime of sensitive data by saving it to the checkpoint.
- ◆ SPARC reduces the lifetime of sensitive data by preventing unintended checkpointing of process-specific memory contents.

## ● Future Work:

- ◆ Detecting and excluding all non-system critical processes that communicate with the process to be excluded.
- ◆ Extend SPARC to exclude confidential **disk information** from being checkpointed.
- ◆ Excluding sensitive data **displayed on the GUI**.

---

**Thank You!**

**Questions?**

---